

tokenmaxxing

Why pick a model when you can win every call?

THE PITCH

Most devs hardcode one AI provider and pay full price — even when a model that costs 20× less would have done the job just as well.

Tokenmaxxing is a single proxy endpoint that routes every AI request to the optimal model based on what you actually care about. Drop in one URL. Stop thinking about providers. Start tokenmaxxing.

ROUTING MODES

- **Cost** Cheapest model per token — GPT-4o Mini at \$0.15/1M.
- **Latency** Fastest model via live rolling latency averages.
- **Quality** Highest-performing model for tasks that demand it.

WHAT'S INCLUDED

- Single proxy API (POST /v1/completion) — drop-in for Anthropic + OpenAI
- Smart routing engine with real-time latency tracking across 6 models
- Credits abstraction layer — users see credits, not provider tokens
- Live dashboard: cost per request, savings vs baseline, provider split
- Interactive architecture diagram (clickable nodes with code examples)
- API playground to test routing decisions in real time · Full OpenAPI spec

THE API

```
POST /v1/completion
{
  "prompt": "Explain the CAP theorem",
  "priority": "cost"    ← cost | latency | quality
}
```

Built with Flask · React · SQLite · Recharts

Mock mode on by default — bring your own Anthropic / OpenAI keys to go live.